

Intro to Using Galaxy

—

For Bioinformatics

Tom Doak

Carrie Ganote

National Center for Genome Analysis Support

September 17, 2013



INDIANA UNIVERSITY



Summary

- Who is NCGAS?
- Galaxy – what is it?
- Galaxy 101 – a guided tour
- Short intro to transcriptome assembly, as an example



INDIANA UNIVERSITY

Who is NCGAS?

The National Center for Genome Analysis Support is based at IU in Bloomington, but caters to a national audience with support from the NSF. We provide computational resources and support for genomics, transcriptomics, and meta projects.



INDIANA UNIVERSITY

Our Services

NCGAS provides support in the form of long- and short-term consultation for genomics, proteomics, transcriptomics, and meta projects. We are happy to answer questions about software, methods, and pipelines; basic Linux use; experimental setup; and interpretation of results.

We administer bioinformatics software installation and upgrades on the Mason cluster at IU, as well as provide access to Mason to users of XSEDE's national infrastructure. We provide support letters for NSF proposals pledging our compute resources.

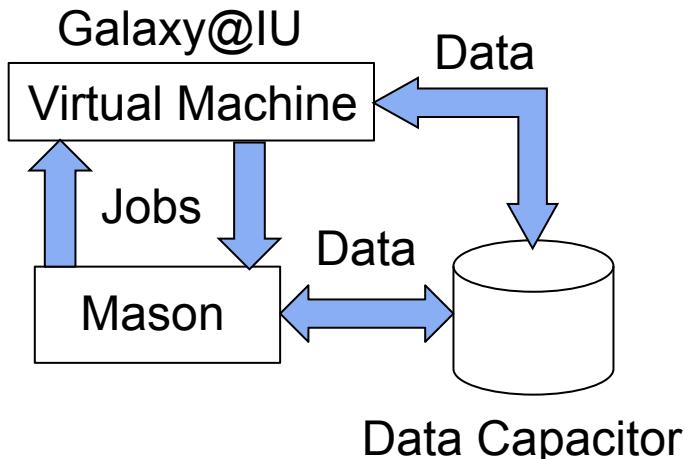
Last, but not least, we install and maintain the local Galaxy instances for Indiana University: IU, NCGAS, and Rockhopper.

What is Galaxy?

Galaxy is a web-based framework for running command-line utilities from a snazzy graphical user interface.

The Galaxy web server that we will be using today is hosted at Indiana University on the XSEDE virtual machines. This is a different “instance” than Galaxy Main, which is hosted at Penn State.

Our instance at IU



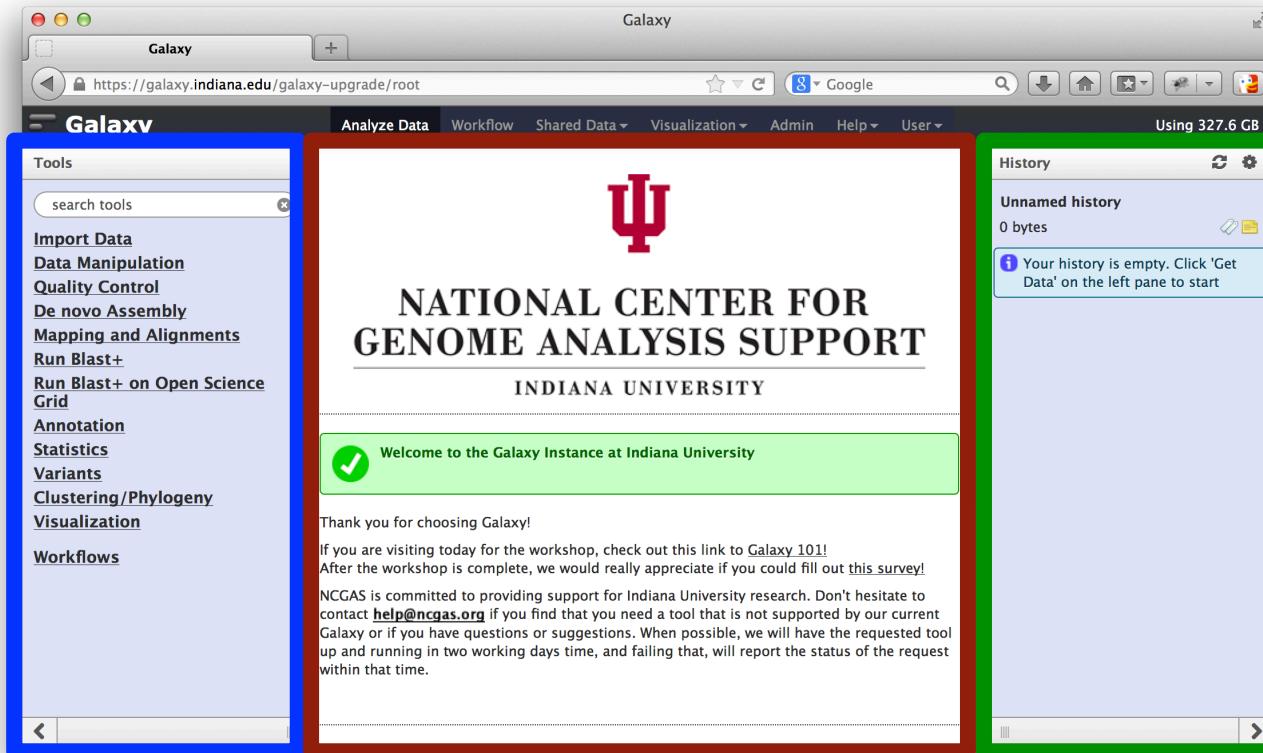
Why choose us:

- IU only – less busy!
- Large RAM jobs possible
- Custom tools on request
- On-site support

Galaxy Main



Galaxy Anatomy and Physiology



Tool bar – contains the available steps to apply to data

History – shows steps previously taken to manipulate input data sets

Focus pane – shows options, parameters, and output for current item.



INDIANA UNIVERSITY

Galaxy 101 – Quick Start

We will depart this slideshow for a short time as we go through the basics of Galaxy using the Galaxy 101 tutorial. You can find a link to it on the home page for galaxy.indiana.edu.

You can choose to follow along either on IU Galaxy or on Galaxy Main – the tool layout is slightly different between the two instances.



INDIANA UNIVERSITY

Today's Menu Item



Cristobal Rojas, La miseria (1886) from Wikipedia.

We will be assembling the DNA Polymerase protein units from the H37Rv strain of *Mycobacterium tuberculosis*, the causative agent of TB, also known as the consumption.

The raw reads originated from the Short Read Archive on NCBI. The accession number for the set is SRX212035.

This dataset consists of paired-end, ~75bp RNA-Seq reads.



INDIANA UNIVERSITY

Let's get some sequence data

Galaxy allows users to publish their data to the entire user base.

A screenshot of the Galaxy web interface. The top navigation bar includes 'Galaxy' (logo), 'Analyze Data', 'Workflow', 'Shared Data' (which is highlighted with a yellow circle), 'Visualization', 'Admin', 'Help', and 'User'. Below the navigation is a search bar ('search tools') and a sidebar with 'Tools' (dropdown), 'Import Data' (selected), and 'Data Manipulation'. The main content area shows a large red Indiana University Psi logo. On the right, there's a 'History' panel with 'Unnamed history' and '0 bytes', and a message 'Your history is empty. Click 'Get''.

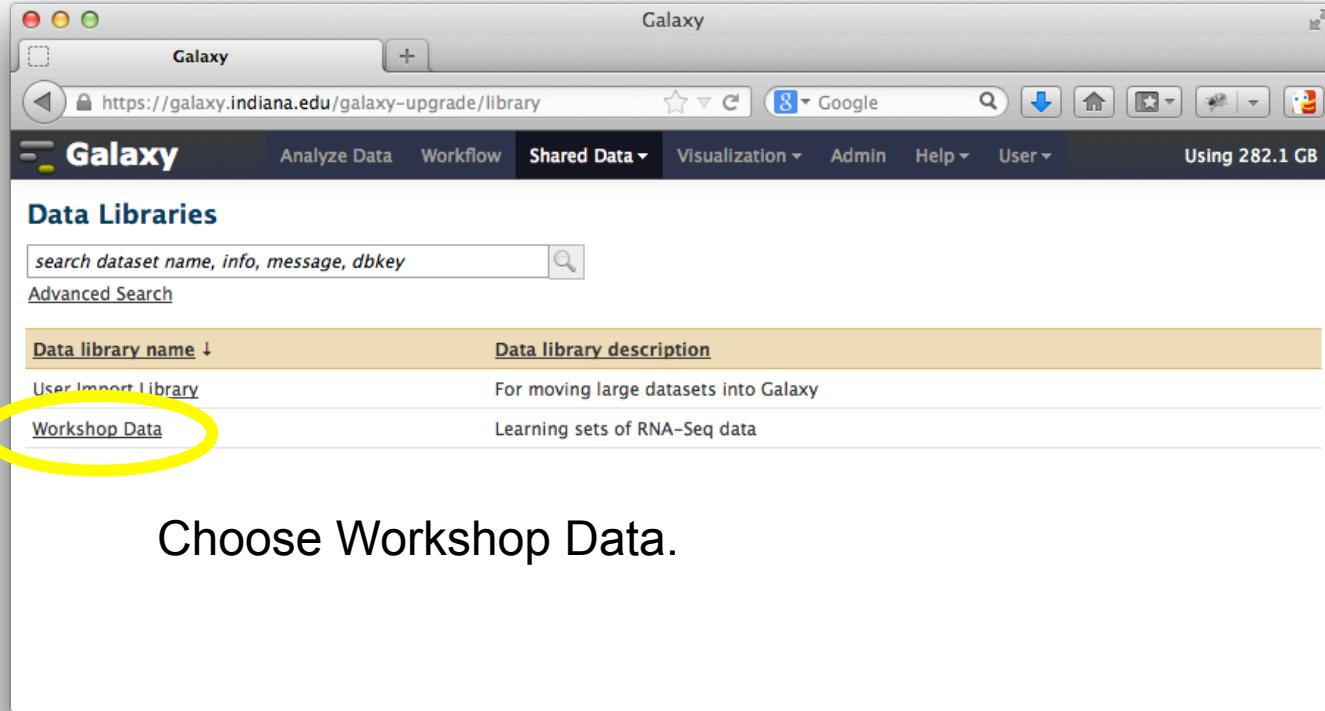
Let's start with "Shared Data" at the top.
Then select Data Libraries from the menu.

A screenshot of the Galaxy web interface, similar to the one above but with a different view. The 'Shared Data' menu is open, showing 'Data Libraries' (which is highlighted with a yellow circle) and other options like 'Published Workflows', 'Published Visualizations', and 'Published Pages'. The bottom of the screen features the 'NATIONAL CENTER FOR GENOME ANALYSIS SUPPORT' logo and 'INDIANA UNIVERSITY'.



INDIANA UNIVERSITY

Let's get some sequence data



The screenshot shows the Galaxy web interface on a Mac OS X system. The title bar says "Galaxy". The address bar shows the URL <https://galaxy.indiana.edu/galaxy-upgrade/library>. The main menu includes "Galaxy", "Analyze Data", "Workflow", "Shared Data", "Visualization", "Admin", "Help", "User", and "Using 282.1 GB". The "Shared Data" menu is currently selected. Below the menu, there is a search bar with placeholder text "search dataset name, info, message, dbkey" and a magnifying glass icon. There is also a link to "Advanced Search". The main content area is titled "Data Libraries". It displays two rows of data:

Data library name ↓	Data library description
User Import Library	For moving large datasets into Galaxy
<u>Workshop Data</u>	Learning sets of RNA-Seq data

A yellow circle highlights the "Workshop Data" row. Below the table, a large text box contains the instruction: "Choose Workshop Data."



INDIANA UNIVERSITY

Let's get some sequence data

Expand folder →

Check both boxes →

Import the Data sets to current history.

Name	Message	Data type	Date uploaded	File size
Galaxy Workshop September '13				
TB_1.fq		fastqsanger	2013-09-15	3.1 MB
TB_2.fq	Right reads	fastqsanger	2013-09-15	3.0 MB

For selected datasets: Import to current history Go

TIP: You can download individual library datasets by selecting "Download this dataset" from the context menu (triangle) next to each dataset's name.

TIP: Several compression options are available for downloading multiple library datasets simultaneously:

- gzip: Recommended for fast network connections
- bzip2: Recommended for slower network connections (smallest file size but takes longer to compress)
- zip: Not recommended but is provided as an option for those who cannot open the above formats



INDIANA UNIVERSITY

Let's get some sequence data

The screenshot shows the Galaxy web interface. At the top, there is a navigation bar with tabs for "Galaxy", "Analyze Data", "Workflow", "Shared Data", "Visualization", "Admin", "Help", "User", and a status message "Using 327.6 GB". A yellow circle highlights the "Analyze Data" tab. Below the navigation bar, the title "Data Library 'Workshop Data'" is displayed. A green banner indicates "2 datasets imported into 1 history: Unnamed history". The main content area shows a table of imported datasets:

Name	Message	Data type	Date uploaded	File size
TB_1.fq		fastqsanger	2013-09-15	3.1 MB
TB_2.fq	Right reads	fastqsanger	2013-09-15	3.0 MB

Below the table, there is a section for "For selected datasets" with a "Import to current history" button and a "Go" button. A tip message states: "TIP: You can download individual library datasets by selecting "Download this dataset" from the context menu (triangle) next to each dataset's name." Another tip message states: "TIP: Several compression options are available for downloading multiple library datasets simultaneously: • gzip: Recommended for fast network connections • bzip2: Recommended for slower network connections (smaller size but takes longer to compress) • zip: Not recommended but is provided as an option for those who cannot open the above formats".

Data set is imported – Click on Analyze Data to return.



INDIANA UNIVERSITY

Step 1: Assess the Quality of Inputs

We will first get an idea of the quality of our input data sets.

The FastQC tool will produce graphical output that makes it easy to gauge the characteristics of the data – quality, patterns, biases, gc content etc.

The screenshot shows the Galaxy web interface. In the top navigation bar, 'Galaxy' is selected. The main content area is titled 'FastQC:Read QC (version 0.51)'. It has several input fields: 'Short read data from your current history:' containing '2: TB_2.fq', 'Title for the output file - to remind you what the job was for:' containing 'FastQC', and 'Contaminant list:' with 'Selection Is Optional'. Below these are 'Execute' and 'Purpose' sections. The 'Purpose' section describes FastQC's function as a quality control tool for raw sequence data. The 'FastQC' section provides a brief overview of the tool's features. On the left sidebar, under 'Quality Control', the 'FastQC:Read QC reports using FastQC' option is highlighted with a yellow circle. The right sidebar shows a 'History' panel with two entries: '2: TB_2.fq' and '1: TB_1.fq'.

Choose either the left or right reads. Compare the results with your neighbor.

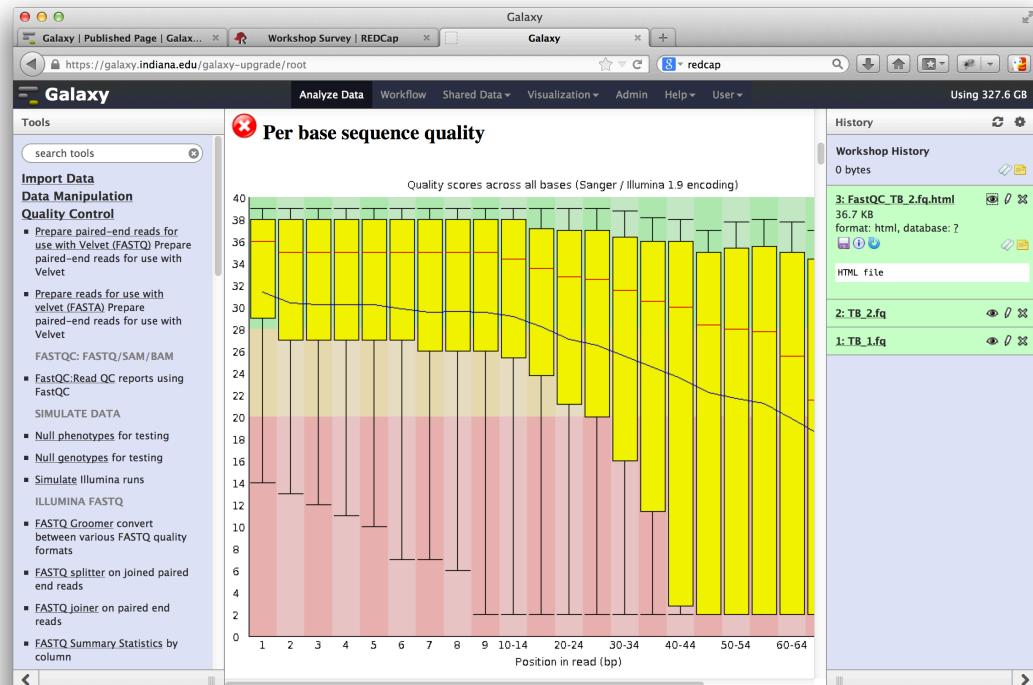


INDIANA UNIVERSITY

Step 1: Assess the Quality of Inputs

The input data usually declines in quality as the reads progress.

The quality score is assigned by the sequencing machine as it reads each base. It is a rough estimate of how ambiguous the signal is.



Sequence: ATGCAT
Quality Score: 39 38 23 19 3 3

Step 2: Trim Input Sequences

We've determined that the input data sets need some work before they are used in downstream processes. We'll use the FASTQ quality trimmer by sliding window to trim reads based on quality score.

The screenshot shows the Galaxy web interface with the title "Galaxy" and the URL "https://galaxy.indiana.edu/galaxy-upgrade/root". The main content area displays the "FASTQ Quality Trimmer (version 1.0.0)" tool. The tool configuration includes:

- FASTQ File: 1: TB_1.fq
- Keep reads with zero length:
- Trim ends: 5' and 3'
- Window size: 1
- Step Size: 1
- Maximum number of bases to exclude from the window during aggregation: 0
- Aggregate action for window: min score
- Trim until aggregate score is: >= 20.0

A yellow circle highlights the "FASTQ Quality Trimmer by sliding window" option in the list of tools on the left sidebar. The right sidebar shows the "History" panel with items: "Workshop History" (898.3 KB), "3: FastQC_TB_2.fq.html" (36.7 KB), "2: TB_2.fq", and "1: TB_1.fq".

Run this tool for both input data sets.



INDIANA UNIVERSITY

Step 3: Rinse, Repeat

Now that the files are trimmed, we will re-assess their quality. If necessary, keep trimming away until you are satisfied with the input files.

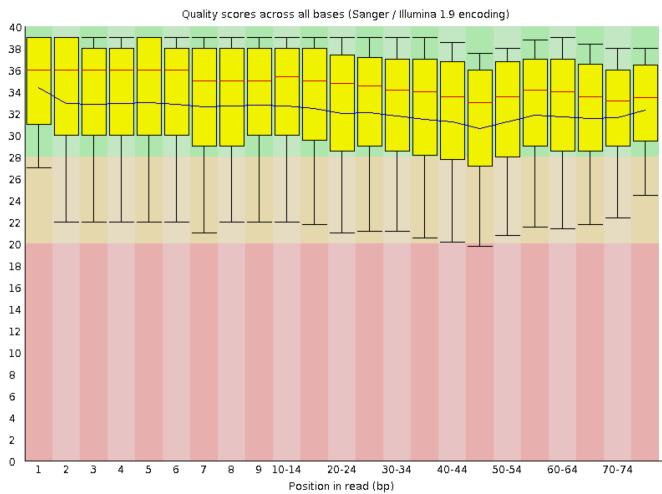
The screenshot shows the Galaxy web interface with the URL <https://galaxy.indiana.edu/galaxy-upgrade/root>. The main panel displays the FastQC tool configuration. The title is "FastQC:Read QC (version 0.51)". The input data is set to "4: Trimmed left reads". The output file title is "FastQC". The contaminant list is set to "Selection is Optional". Below the form, there is a "Purpose" section describing FastQC's function as a quality control tool for raw sequence data. The "FastQC" section notes it is a Galaxy wrapper for the external package. The "Summary" section shows two history items: "5: Trimmed right reads" and "4: Trimmed left reads", each listing sequence details like length and database format.

I renamed my trimmed files to help me keep them straight.

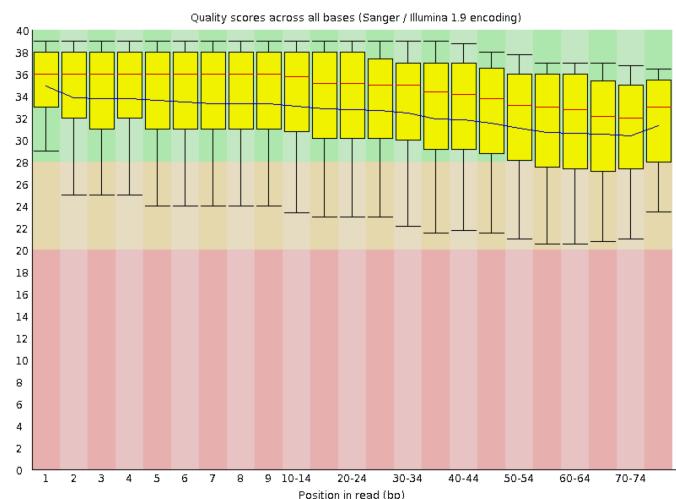
Step 3: Rinse, Repeat

Pictured are the left and right reads after trimming is complete.
These will do!

✓ Per base sequence quality



✓ Per base sequence quality



Step 4: Assembly

Next we will put the reads together to create a complete picture of the actively transcribed genes of the sample organism.

Trinity is a *de novo* assembler that has been optimized for use on Mason. We will use it to assemble our reads.

The screenshot shows the Galaxy web interface. On the left, a sidebar lists various tools: Import Data, Data Manipulation, Quality Control, De novo Assembly, Trimming and Alignment, Run Blast+, Run Blast+ on Open Science Grid, Annotation, Statistics, Variants, Clustering/Phylogeny, Visualization, and Workflows. The 'De novo Assembly' section is highlighted with a yellow circle. The main workspace shows the 'Trinity - Executes on Mason (version 0.0.1)' tool configuration. It has dropdown menus for 'Paired or Single-end data?' (set to 'Paired'), 'Left/Forward strand reads' (set to '4: Trimmed left reads'), and 'Right/Reverse strand reads' (set to '5: Trimmed right reads'). Below these are fields for 'Strand-specific Library Type' (set to 'None'), 'Paired Fragment Length' (set to '300'), and 'Maximum length expected between fragment pairs'. There are also dropdowns for 'Is it strand specific data?' (set to 'No') and 'Use Additional Params?' (set to 'No'). A 'How long will your job need?' dropdown is set to '1 hr'. At the bottom is a blue 'Execute' button. A note below the execute button says: 'If you have no idea about the walltime, contact the Galaxy team for assistance.' To the right of the tool configuration is a 'History' panel showing the workflow's progress. It includes sections for 'Workshop History' (7.5 MB), 'FastQC_Trimmed right reads.html' (0 reads), 'FastQC_Trimmed left reads.html' (0 reads), 'Trimmed right reads.html' (12,913 sequences, 14,024 FASTQ reads processed, 11,111 reads of zero length excluded), and two sequence snippets: one for 'SR638895.6844' and another for 'SR638895.6859'. Below these are sections for 'Trimmed left reads.html' (13,087 sequences, 14,117 FASTQ reads processed, 1,030 reads of zero length excluded) and 'Trimmed right reads.html' (13,087 sequences, 14,117 FASTQ reads processed, 1,030 reads of zero length excluded).



INDIANA UNIVERSITY

It finished! We're done, right?

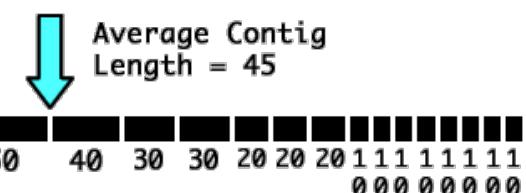
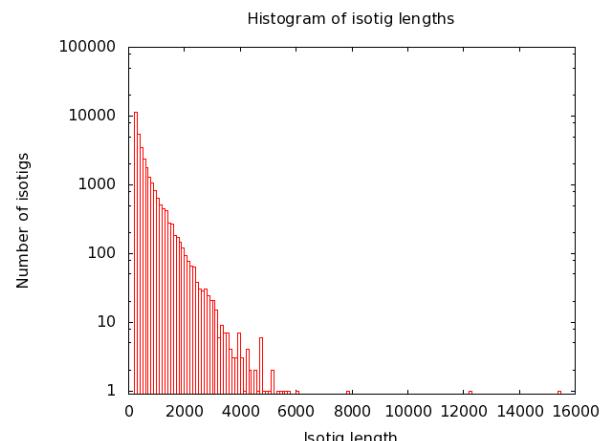
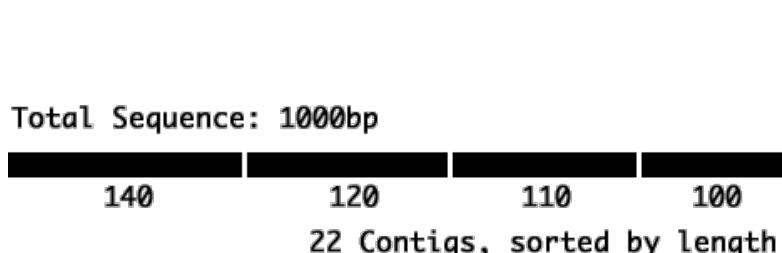
An assembler solves a computer problem of putting together a puzzle from tiny pieces. The output of the assembler is a guess – but we don't know how accurate it is. We could look at:

- Basic stats of the assembly – “Contigs”
 - Number of “Contigs” vs. Expected Number
 - N50 – a weighted average
 - Average Length
 - Max Length
- Check contigs against known genes with Blast (large or rare transcripts)

Step 5: Assessing Quality of Assembly

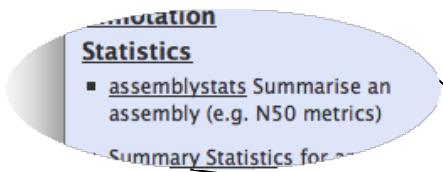
Important statistics for assembly quality:
Contig Length Distribution

Assemblies will typically produce a number of complete contigs representing whole transcripts, and a large number of partial transcripts. This biases the average contig length toward the low end. The N50 is a measure weighted by total sequence length in the assembly.



Step 5: Assessing Quality of Assembly

Getting these stats in Galaxy:



Run assemblystats to get a summary and histograms of your contig length distribution.

A screenshot of the Galaxy assemblystats tool configuration page. The tool version is 1.0.1. The 'Type of read:' dropdown is set to 'Isotig (if from transcriptomic assembly)'. There is a checkbox for 'Output histogram with bin sizes=1' which is unchecked. The 'Source file in FASTA format:' dropdown contains the value '84: Trinity on data 20 and data 21: Assembled Transcripts'. At the bottom, there is an 'Execute' button. To the right, the 'History' panel shows two entries: '240: Sorted contigs' containing a FASTA sequence, and '239: Assembly statistics' containing a table of assembly statistics including N50, contig count, and total length.

Step 6: Check Against Database

For this last step, we'll check to see how well our assembled transcripts compare to what we already know.

Use this step to give a rough annotation of genes, to make sure that your transcripts are from nuclear genes, or to gauge how complete your sequence is.

The screenshot shows the Galaxy web interface. On the left, a sidebar titled 'Tools' lists several bioinformatics tools, with 'Filter sequences by length' highlighted by a yellow circle. The main panel displays the 'Filter sequences by length (version 1.1)' tool configuration. It includes fields for 'Minimal length' (set to 0) and 'Maximum length' (set to 555), a 'Execute' button, and a tip about setting minimal length to 0. Below these, a section titled 'What it does' describes the tool's function: 'Outputs sequences between Minimal length and Maximum length.' The 'History' panel on the right shows a list of completed steps, including 'Workshop History' and various assembly and transcript-related tasks.

For sake of time, we'll just Blast one gene. Filter out to get the smallest.

Step 6: Check Against Database

We will use Blastx to search the NR database for our gene.

Use default search settings for this test set.

The screenshot shows the Galaxy web interface at <https://galaxy.indiana.edu/galaxy-upgrade/root>. The main panel displays the 'NCBI BLAST+ blastx (version 0.0.17)' tool configuration. The 'Tools' sidebar on the left lists several BLAST-related tools, with 'NCBI BLAST+ blastx' highlighted by a yellow circle. The configuration form includes fields for 'Nucleotide query sequence(s)', 'Subject database/sequences', 'Protein BLAST database', 'Query genetic code', 'Set expectation value cutoff', 'Output format', and 'Advanced Options'. A note at the bottom cautions that database searches may take time. The right side shows a 'History' panel with a list of previous workflows and their outputs.

Make sure to choose Pairwise HTML output for readability.

Step 6: Check Against Database

We see the expected genes as the top hits!

The screenshot shows a Galaxy web interface window. The title bar says "Galaxy" and the address bar shows the URL "https://galaxy.indiana.edu/galaxy-upgrade/root". The main content area displays a BLAST search results table and some sequence alignments.

Tools:

- nucleotide database with nucleotide query sequence(s)
- NCBI BLAST+ tblastx Search translated nucleotide database with translated nucleotide query sequence(s)
- NCBI BLAST+ blastx Search protein database with translated nucleotide query sequence(s)
- NCBI BLAST+ database info Show BLAST database information from blastdbcmd
- Megablast compare short reads against htgs, nt, and wgs databases
- Parse blast XML output

Analyze Data:

ref	WP_009744290.1	DNA polymerase III subunit delta' [Actinomyc...
ref	WP_005963184.1	DNA-directed DNA polymerase III subunit delta...
ref	WP_010525702.1	DNA polymerase III subunit delta' [Nesterenk...
ref	WP_017201448.1	hypothetical protein [Microbacterium barkeri]
ref	YP_830214.11	DNA polymerase III subunit delta' [Arthrobacter...
ref	WP_019482356.1	hypothetical protein [Arthrobacter sp. TB 23]
ref	WP_004806969.1	DNA polymerase III subunit delta' [Actinomyc...

Sequence Alignments:

```
>ref|WP_003907097.1| DNA polymerase III subunit delta, partial [Mycoba  
gb|EFD75343.1| DNA polymerase III subunit delta [Mycobacterium tuberc  
GM 1503]  
Length=354  
  
Score = 246 bits (627), Expect = 3e-77  
Identities = 181/181 (100%), Positives = 181/181 (100%), Gaps = 0/181  
Frame = -3  
  
Query 553 TDPQARQRERALGLARDAAATPSRAYAAAEEELVAGAAEAEALALTAQRIEATEELRTA  
Sbjct 174 TDPQARQRERALGLARDAAATPSRAYAAAEEELVAGAAEAEALALTAQRIEATEELRTA  
  
Query 373 aggtgkgtgaalrgatqAMKDLERRQKSQTRASRDLALRALIDLATYFRDALLVAAH  
Sbjct 234 AGGTGKGTGAALRGATGAMKDLERRQKSQTRASRDLALRALIDLATYFRDALLVAAH  
  
Query 193 GVRANHPDMADRAVAALAAHAPPERLLRCIEAVLACREALAVNVKPKFAVDAMVATIGC  
Sbjct 294 GVRANHPDMADRAVAALAAHAPPERLLRCIEAVLACREALAVNVKPKFAVDAMVATIGC  
  
Query 13 R 11
```

History:

- 17: blastx on db 484.1 KB format: html, database: ?
- 16: Filter sequences by length on data 9 1 sequences format: fasta, database: ?

Output:

```
>comp7_c0_seq1 len=555 path=[1:0-55:  
GCCGTCCCTACCCAGTTCTGGCGATGGTGGCG  
GTTTGACGTTGACCGCTAGCGCTTCCCTGCACGCC  
GCCGCTCCGGCGGGCGTGGCGGGCCAGCGCAGCA
```

We could limit the number of hits depending on output desired.



Step 7..?

RNA-Seq is a very versatile technology. You can use the data for:

- Gene discovery based on transcripts
- Genome evidence – introns, exons, junction
- Gene expression patterns
- SNP calling/other variants
- Protein divergence between samples

We have gotten to the assembly step, but there is a lot to learn about the data now that it is put together. A foundation in the use of Galaxy coupled with Indiana University resources will enable you to reach these goals.



INDIANA UNIVERSITY

Fin

Thanks for watching!
Questions and comments:
Email help@ncgas.org